

```
In [57]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn import svm
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
```

Data Set Up

Welcome to my final project assignment for COGS118A.

```
In [58]: bank_df = pd.read_csv("Bank/bank-full.csv", sep = ";")
display(bank_df)
```

| | age | job | marital | education | default | balance | housing | loan | contact | day | month | duration | campaign | pdays |
|-------|-----|--------------|----------|-----------|---------|---------|---------|------|-----------|-----|-------|----------|----------|-------|
| 0 | 58 | management | married | tertiary | no | 2143 | yes | no | unknown | 5 | may | 261 | 1 | -1 |
| 1 | 44 | technician | single | secondary | no | 29 | yes | no | unknown | 5 | may | 151 | 1 | -1 |
| 2 | 33 | entrepreneur | married | secondary | no | 2 | yes | yes | unknown | 5 | may | 76 | 1 | -1 |
| 3 | 47 | blue-collar | married | unknown | no | 1506 | yes | no | unknown | 5 | may | 92 | 1 | -1 |
| 4 | 33 | unknown | single | unknown | no | 1 | no | no | unknown | 5 | may | 198 | 1 | -1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 45206 | 51 | technician | married | tertiary | no | 825 | no | no | cellular | 17 | nov | 977 | 3 | -1 |
| 45207 | 71 | retired | divorced | primary | no | 1729 | no | no | cellular | 17 | nov | 456 | 2 | -1 |
| 45208 | 72 | retired | married | secondary | no | 5715 | no | no | cellular | 17 | nov | 1127 | 5 | 184 |
| 45209 | 57 | blue-collar | married | secondary | no | 668 | no | no | telephone | 17 | nov | 508 | 4 | -1 |
| 45210 | 37 | entrepreneur | married | secondary | no | 2971 | no | no | cellular | 17 | nov | 361 | 2 | 188 |

```
In [59]: mushroom_df = pd.read_csv("Mushroom/agaricus-lepiota.data")
display(mushroom_df)
```

| | p | x | s | n | t | p.1 | f | c | n.1 | k | ... | s.2 | w | w.1 | p.2 | w.2 | o | p.3 | k.1 | s.3 | u |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | e | x | s | y | t | a | f | c | b | k | ... | s | w | w | p | w | o | p | n | n | g |
| 1 | e | b | s | w | t | l | f | c | b | n | ... | s | w | w | p | w | o | p | n | n | m |
| 2 | p | x | y | w | t | p | f | c | n | n | ... | s | w | w | p | w | o | p | k | s | u |
| 3 | e | x | s | g | f | n | f | w | b | k | ... | s | w | w | p | w | o | e | n | a | g |
| 4 | e | x | y | y | t | a | f | c | b | n | ... | s | w | w | p | w | o | p | k | n | g |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 8118 | e | k | s | n | f | n | a | c | b | y | ... | s | o | o | p | o | o | p | b | c | l |
| 8119 | e | x | s | n | f | n | a | c | b | y | ... | s | o | o | p | n | o | p | b | v | l |
| 8120 | e | f | s | n | f | n | a | c | b | n | ... | s | o | o | p | o | o | p | b | c | l |
| 8121 | p | k | y | n | f | y | f | c | n | b | ... | k | w | w | p | w | o | e | w | v | l |
| 8122 | e | x | s | n | f | n | a | c | b | y | ... | s | o | o | p | o | o | p | o | c | l |

8123 rows × 23 columns

```
In [60]: income_df = pd.read_csv("Income/adult.data", delimiter=",")
display(income_df)
```

| | 39 | State-gov | 77516 | Bachelors | 13 | Never-married | Adm-clerical | Not-in-family | White | Male | 2174 | 0 | 40 | United-States | <=50K |
|-------|-----|------------------|--------|------------|-----|--------------------|-------------------|---------------|-------|--------|------|-----|-----|---------------|-------|
| 0 | 50 | Self-emp-not-inc | 83311 | Bachelors | 13 | Married-civ-spouse | Exec-managerial | Husband | White | Male | 0 | 0 | 13 | United-States | <=50K |
| 1 | 38 | Private | 215646 | HS-grad | 9 | Divorced | Handlers-cleaners | Not-in-family | White | Male | 0 | 0 | 40 | United-States | <=50K |
| 2 | 53 | Private | 234721 | 11th | 7 | Married-civ-spouse | Handlers-cleaners | Husband | Black | Male | 0 | 0 | 40 | United-States | <=50K |
| 3 | 28 | Private | 338409 | Bachelors | 13 | Married-civ-spouse | Prof-specialty | Wife | Black | Female | 0 | 0 | 40 | Cuba | <=50K |
| 4 | 37 | Private | 284582 | Masters | 14 | Married-civ-spouse | Exec-managerial | Wife | White | Female | 0 | 0 | 40 | United-States | <=50K |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 32555 | 27 | Private | 257302 | Assoc-acdm | 12 | Married-civ-spouse | Tech-support | Wife | White | Female | 0 | 0 | 38 | United-States | <=50K |

One-Hot Encoding

While our data looks nice and formatted, we sadly cannot work with categorical labels to perform our model training. To resolve this, we utilize one-hot encoding to transform the set of each possible nominal values from every column into a new column span that treats each label as a 0 or 1. For other columns with numerical values, we simply treat their values the same. Thankfully, numpy comes with a method to do this for us automatically for each of our datasets

```
In [61]: bank_encoded = pd.get_dummies(bank_df)
bank_encoded = bank_encoded.replace({True: 1, False: 0})

mushroom_encoded = pd.get_dummies(mushroom_df)
mushroom_encoded = mushroom_encoded.replace({True: 1, False: 0})

income_encoded = pd.get_dummies(income_df)
income_encoded = income_encoded.replace({True: 1, False: 0})
```

```
In [62]: display(bank_encoded)
```

| | age | balance | day | duration | campaign | pdays | previous | job_admin. | job_blue-collar | job_entrepreneur | ... | month_may | month_ |
|-------|-----|---------|-----|----------|----------|-------|----------|------------|-----------------|------------------|-----|-----------|--------|
| 0 | 58 | 2143 | 5 | 261 | 1 | -1 | 0 | 0 | 0 | 0 | ... | 1 | |
| 1 | 44 | 29 | 5 | 151 | 1 | -1 | 0 | 0 | 0 | 0 | ... | 1 | |
| 2 | 33 | 2 | 5 | 76 | 1 | -1 | 0 | 0 | 0 | 1 | ... | 1 | |
| 3 | 47 | 1506 | 5 | 92 | 1 | -1 | 0 | 0 | 1 | 0 | ... | 1 | |
| 4 | 33 | 1 | 5 | 198 | 1 | -1 | 0 | 0 | 0 | 0 | ... | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 45206 | 51 | 825 | 17 | 977 | 3 | -1 | 0 | 0 | 0 | 0 | ... | 0 | |
| 45207 | 71 | 1729 | 17 | 456 | 2 | -1 | 0 | 0 | 0 | 0 | ... | 0 | |
| 45208 | 72 | 5715 | 17 | 1127 | 5 | 184 | 3 | 0 | 0 | 0 | ... | 0 | |
| 45209 | 57 | 668 | 17 | 508 | 4 | -1 | 0 | 0 | 1 | 0 | ... | 0 | |

```
In [63]: display(mushroom_encoded)
```

| | p_e | p_p | x_b | x_c | x_f | x_k | x_s | x_x | s_f | s_g | ... | s.3_s | s.3_v | s.3_y | u_d | u_g | u_l | u_m | u_p | u_u | u_w |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-------|-------|-------|-----|-----|-----|-----|-----|-----|-----|
| 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| 1 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 |
| 2 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | ... | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 |
| 3 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| 4 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 8118 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 8119 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | ... | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 8120 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 8121 | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | ... | 0 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| 8122 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |

8123 rows × 119 columns

```
In [64]: display(income_encoded)
```

| | 39 | 77516 | 13 | 2174 | 0 | 40 | State-gov_? Federal-gov | State-gov_ Local-gov | State-gov_ Never-worked | ... | United-States_ Scotland | United-States_ South | United-States_ Taiwan | United-States_ Thailand | United-States_ Trinidad&Tobago |
|-------|-----|--------|-----|-------|-----|-----|----------------------------|-------------------------|----------------------------|-----|----------------------------|-------------------------|--------------------------|----------------------------|-----------------------------------|
| 0 | 50 | 83311 | 13 | 0 | 0 | 13 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| 1 | 38 | 215646 | 9 | 0 | 0 | 40 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| 2 | 53 | 234721 | 7 | 0 | 0 | 40 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| 3 | 28 | 338409 | 13 | 0 | 0 | 40 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| 4 | 37 | 284582 | 14 | 0 | 0 | 40 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 32555 | 27 | 257302 | 12 | 0 | 0 | 38 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| 32556 | 40 | 154374 | 9 | 0 | 0 | 40 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| 32557 | 58 | 151910 | 9 | 0 | 0 | 40 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| 32558 | 22 | 201490 | 9 | 0 | 0 | 20 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |
| 32559 | 52 | 287927 | 9 | 15024 | 0 | 40 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 | 0 |

32560 rows × 110 columns

Converting to Matrices

Before we rush ahead, we need to convert each dataframe into a form that we can easily work with so we can utilize complex mathematical operations. We'll extract each value in the dataframes and transfer them over to a numpy array

```
In [65]: bank = bank_encoded.values
mushroom = mushroom_encoded.values
income = income_encoded.values

# Verifying that the shape matches their dataframe shape
print(bank.shape)
print(mushroom.shape)
print(income.shape)
```

```
(45211, 53)
(8123, 119)
(32560, 110)
```

Cleaning Up the Data

With the data mostly done, we'll focus on setting up the classifiers. We also have to address the fact that we can't keep the columns that classify our observations since they are not features. Since we already have each table in matrix form, we can easily remove and extract our labels for each data set! Additionally, we must shuffle the data.

```
In [66]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
```

```
In [67]: # Extract the corresponding column for the labels, then remove the last two columns
bank_labels = (bank[:, -1]).reshape(-1, 1).astype(float)
bank = bank[:, :-2]

mushroom_labels = (mushroom[:, 1]).reshape(-1, 1).astype(float)
mushroom = mushroom[:, 2:]

income_labels = (income[:, -1]).reshape(-1, 1).astype(float)
income = income[:, :-2]

print('Bank Shape: {}'.format(bank.shape))
print('Mushroom Shape: {}'.format(mushroom.shape))
print('Income Shape: {}'.format(income.shape))

# print(bank_labels.shape)
# print(mushroom_labels.shape)
# print(income_labels.shape)

# Convert every false (0) to -1 in our labels arrays
unique, counts = np.unique(mushroom_labels, return_counts=True)
count_dict = dict(zip(unique, counts))
print(count_dict) # {0: 7, 1: 4, 2: 1, 3: 2, 4: 1}

bank_labels[bank_labels == 0] = -1
mushroom_labels[mushroom_labels == 0] = -1
income_labels[income_labels == 0] = -1

# print(bank_labels)
# print(mushroom_labels)
# print(income_labels)

# Stack the labels with their original tables to shuffle
bank = np.hstack((bank, bank_labels))
mushroom = np.hstack((mushroom, mushroom_labels))
income = np.hstack((income, income_labels))

np.random.seed(1)
np.random.shuffle(bank)
np.random.shuffle(mushroom)
np.random.shuffle(income)
```

```
Bank Shape: (45211, 51)
Mushroom Shape: (8123, 117)
Income Shape: (32560, 108)
{0.0: 4208, 1.0: 3915}
```

```

In [68]: # For class weights
bank_unique, bank_counts = np.unique(bank_labels, return_counts=True)
bank_count_dict = dict(zip(bank_unique, bank_counts))

income_unique, income_counts = np.unique(income_labels, return_counts=True)
income_count_dict = dict(zip(income_unique, income_counts))

mushroom_unique, mushroom_counts = np.unique(mushroom_labels, return_counts=True)
mushroom_count_dict = dict(zip(mushroom_unique, mushroom_counts))

# Extract features and labels for calculating proportions of each class
bank_x = bank[:, :-1]
bank_y = bank[:, -1]

income_x = income[:, :-1]
income_y = income[:, -1]

mushroom_x = mushroom[:, :-1]
mushroom_y = mushroom[:, -1]

# Create a dict of counts of each class for each dataset
bank_weight_1_neg = len(bank_x)/bank_count_dict.get(-1)
bank_weight_1_pos = len(bank_x)/bank_count_dict.get(1)
bank_class_weights = {-1: bank_weight_1_neg , 1: bank_weight_1_pos }

income_weight_1_neg = len(income_x)/income_count_dict.get(-1)
income_weight_1_pos = len(income_x)/income_count_dict.get(1)
income_class_weights = {-1: income_weight_1_neg , 1: income_weight_1_pos }

mushroom_weight_1_neg = len(mushroom_x)/mushroom_count_dict.get(-1)
mushroom_weight_1_pos = len(mushroom_x)/mushroom_count_dict.get(1)
mushroom_class_weights = {-1: mushroom_weight_1_neg , 1: mushroom_weight_1_pos }

```

Support Vector Machines

We won't use the default SVM library but rather SVCLinear. It is similar to SVC with parameter kernel='linear', but implemented in terms of liblinear rather than libsvm, so it has more flexibility in the choice of penalties and loss functions and should scale better to large numbers of samples. The main differences between LinearSVC and SVC lie in the loss function used by default, and in the handling of intercept regularization between those two implementations.

The goal here is to find the best hyperparameter C

```

In [69]: from sklearn.svm import LinearSVC
from sklearn.metrics import classification_report
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import accuracy_score
import scipy.io as sio
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

```

```

In [70]: # Hyperparameter List
C_list = [0.1, 1, 10, 100, 1000, 10000]

```

```
In [71]: # Draw heatmaps for result of grid search.
def draw_heatmap(errors, param_list, title):
    plt.figure(figsize = (2,4))
    ax = sns.heatmap(errors, annot=True, fmt='.4f', yticklabels=param_list, xticklabels=[])
    ax.collections[0].colorbar.set_label('error')
    ax.set(ylabel='hyper parameter')
    bottom, top = ax.get_ylim()
    ax.set_ylim(bottom + 0.5, top - 0.5)
    plt.title(title)
    plt.show()
```

```
In [72]: def calcSVMetrics(X_train,X_test, Y_train,Y_test, C_List, class_weights):

    clf = LinearSVC(dual=False, class_weight=class_weights)
    param_grid = {'C': C_list}

    # Perform 3-Fold cross validation for each hyperparameter
    grid_search = GridSearchCV(clf, param_grid, cv=3, return_train_score=True )

    # Fit the model
    grid_search.fit(X_train, Y_train)

    # Gather the results
    opt_C = grid_search.best_params_['C']

    cross_validation_accuracies = grid_search.cv_results_['mean_test_score']
    cross_validation_errors = 1 - cross_validation_accuracies.reshape(-1,1)

    mean_training_accuracies = grid_search.cv_results_['mean_train_score']
    mean_training_errors = 1 - mean_training_accuracies.reshape(-1,1)

    Y_pred = grid_search.best_estimator_.predict(X_test)
    test_accuracy = accuracy_score(Y_pred, Y_test)
    test_error = 1 - sum(Y_pred == Y_test) / len(X_test)

    return opt_C,cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_train
```

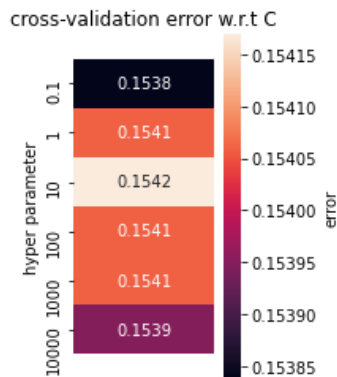
SVM for BANK

```
In [79]: # Array to track accuracies for each partition
svm_bank_accs = []

# 20% Training and 80% Testing
bank_x_train_20, bank_x_test_80, bank_y_train_20, bank_y_test_80 = train_test_split(bank_x, bank_y, test_size=0.2,
opt_C, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accuracies))

svm_bank_accs.append(test_accuracy)
```



Best C: 0.1

Test error: 0.15969476623627965

Test accuracy: 0.8403052337637203

Avg training error per C: [[0.14902676]

[0.14924795]

[0.14891617]

[0.14875028]


[0.14897147]

[0.14869498]]

Avg training accuracies per C: [0.85097324 0.85075205 0.85108383 0.85124972 0.85102853 0.85130502]

```
In [80]: # 50% Training and 50% Testing
bank_x_train_50, bank_x_test_50, bank_y_train_50, bank_y_test_50 = train_test_split(bank_x, bank_y, test_size=0.5, random_state=0)
opt_C, cross_validation_accs, cross_validation_errors, mean_training_accs, mean_training_errors = grid_search_svm()

draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accs))
svm_bank_accs.append(test_accuracy)
```




```
Best C: 10
Test error: 0.14987171547376799
Test accuracy: 0.850128284526232
Avg training error per C: [[0.14353019]
 [0.14350807]
 [0.14361867]
 [0.14346384]
 [0.14346384]
 [0.14359655]]
Avg training accuracies per C: [0.85646981 0.85649193 0.85638133 0.85653616 0.85653616 0.85640345]
```



```
In [81]: # 80% Training and 20% Testing
bank_x_train_80, bank_x_test_20, bank_y_train_80, bank_y_test_20 = train_test_split(bank_x, bank_y, test_size=0.2, random_state=42)
opt_C, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors = grid_search_svm()

draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accuracies))

svm_bank_accs.append(test_accuracy)
```




```
Best C: 0.1
Test error: 0.15348888643149394
Test accuracy: 0.8465111135685061
Avg training error per C: [[0.15011889]
 [0.15017419]
 [0.15034008]
 [0.15039538]
 [0.15011889]
 [0.15028478]]
Avg training accuracies per C: [0.84988111 0.84982581 0.84965992 0.84960462 0.84988111 0.84971522]
```

SVM For INCOME

```
In [85]: svm_income_accs = []
# 20% Training and 80% Testing
income_x_train_20, income_x_test_80, income_y_train_20, income_y_test_80 = train_test_split(income_x, income_y, test_size=0.2, random_state=0)
opt_C, cross_validation_accs, cross_validation_errors, mean_training_accs, mean_training_errors = grid_search(svm_model, income_x_train_20, income_y_train_20, income_x_test_80, income_y_test_80)

draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accs))

svm_income_accs.append(test_accuracy)
```



Best C: 1000

Test error: 0.1992859336609336

Test accuracy: 0.8007140663390664

Avg training error per C: [[0.19878695]

[0.20845929]

[0.19794204]

[0.20124251]

[0.19172213]


[0.20170371]]

Avg training accuracies per C: [0.80121305 0.79154071 0.80205796 0.79875749 0.80827787 0.79829629]

```
In [86]: # 50% Training and 50% Testing
income_x_train_50, income_x_test_50, income_y_train_50, income_y_test_50 = train_test_split(income_x, income_y,
opt_C, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accuracies))

svm_income_accs.append(test_accuracy)
```

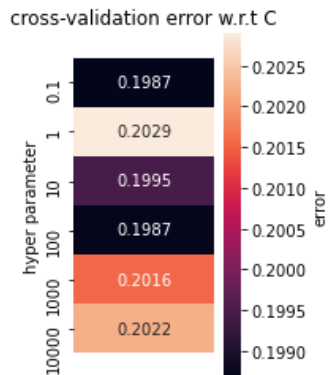


```
Best C: 0.1
Test error: 0.18906633906633907
Test accuracy: 0.8109336609336609
Avg training error per C: [[0.19367318]
[0.1941646 ]
[0.19702097]
[0.19490163]
[0.19656026]
[0.19241396]]
Avg training accuracies per C: [0.80632682 0.8058354 0.80297903 0.80509837 0.80343974 0.80758604]
```

```
In [87]: # 80% Training and 20% Testing
income_x_train_80, income_x_test_20, income_y_train_80, income_y_test_20 = train_test_split(income_x, income_y,
opt_C, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accuracies))

svm_income_accs.append(test_accuracy)
```



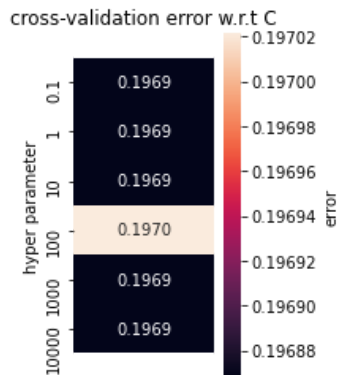
```
Best C: 0.1
Test error: 0.1904176904176904
Test accuracy: 0.8095823095823096
Avg training error per C: [[0.19890202]
[0.20199257]
[0.2000922 ]
[0.19859488]
[0.20155109]
[0.20160863]]
Avg training accuracies per C: [0.80109798 0.79800743 0.7999078 0.80140512 0.79844891 0.79839137]
```

SVM For MUSHROOM

```
In [88]: svm_mushroom_accs = []

# 20% Training and 80% Testing
mushroom_x_train_20, mushroom_x_test_80, mushroom_y_train_20, mushroom_y_test_80 = train_test_split(income,
opt_C, cross_validation_accuaries, cross_validation_errors, mean_training_accuaries, mean_training_errors)


draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accuaries))
svm_mushroom_accs.append(test_accuracy)
```



```
Best C: 0.1
Test error: 0.2021652334152334
Test accuracy: 0.7978347665847666
Avg training error per C: [[0.19702089]
[0.19702089]
[0.19702089]
[0.19702089]
[0.19702089]]
Avg training accuracies per C: [0.80297911 0.80297911 0.80297911 0.80297911 0.80297911 0.80297911]
```

```
In [89]: # 50% Training and 50% Testing
mushroom_x_train_50, mushroom_x_test_50, mushroom_y_train_50, mushroom_y_test_50 = train_test_split(income,
opt_C, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)


draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accuracies))
svm_mushroom_accs.append(test_accuracy)
```



```
Best C: 0.1
Test error: 0.20399262899262904
Test accuracy: 0.796007371007371
Avg training error per C: [[0.20042998]
[0.20049141]
[0.20021501]
[0.20042998]
[0.20042998]
[0.20042998]]
Avg training accuracies per C: [0.79957002 0.79950859 0.79978499 0.79957002 0.79957002 0.79957002]
```

```
In [90]: # 80% Training and 20% Testing
mushroom_x_train_80, mushroom_x_test_20, mushroom_y_train_80, mushroom_y_test_20 = train_test_split(income,
opt_C, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_error)

draw_heatmap(cross_validation_errors, C_list, title='cross-validation error w.r.t C')
print("Best C: {}".format(opt_C))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per C: {}".format(mean_training_errors))
print("Avg training accuracies per C: {}".format(mean_training_accuracies))
svm_mushroom_accs.append(test_accuracy)
```



```
Best C: 1
Test error: 0.2010135135135135
Test accuracy: 0.7989864864864865
Avg training error per C: [[0.20181971]
 [0.2018389 ]
 [0.20166614]
 [0.20168533]
 [0.20155097]
 [0.20181971]]
Avg training accuracies per C: [0.79818029 0.7981611  0.79833386 0.79831467 0.79844903 0.79818029]
```

SVM Results

Now that we've performed training and testing on all three datasets and partitions, let's average all the test accuracies. We'll use this to help with our comparison against Caruana's findings.

```
In [91]: average_accuracy = np.sum([a + b + c for a, b, c in zip(svm_bank_accs, svm_income_accs, svm_mushroom_accs)])
print(svm_bank_accs)
print(svm_income_accs)
print(svm_mushroom_accs)
print("Average SVM accuracy {}".format(average_accuracy))

[0.8403052337637203, 0.850128284526232, 0.8465111135685061]
[0.8007140663390664, 0.8109336609336609, 0.8095823095823096]
[0.7978347665847666, 0.796007371007371, 0.7989864864864865]
Average SVM accuracy 0.8167781436435688
```

K Nearest Neighbors

```
In [92]: import scipy
from matplotlib.colors import ListedColormap
from functools import partial
from sklearn.neighbors import KNeighborsClassifier

# Hyperparameter List of possible K's
# Because of the scope of the project, I capped it at 15 since KNN involves expensive operations
k_range = list(range(1, 16))
```

```
In [93]: def calcKNNMetrics(X_train,X_test, Y_train,Y_test,k_range):

    param_grid = dict(n_neighbors=k_range)
    clf = KNeighborsClassifier(algorithm = 'kd_tree', weights='distance')

    grid_search = GridSearchCV(clf, param_grid, cv=3, return_train_score=True,verbose=1, )

    # Fit the model
    grid_search.fit(X_train, Y_train)

    # Gather the results
    opt_K = grid_search.best_params_

    cross_validation_accuracies = grid_search.cv_results_['mean_test_score']
    cross_validation_errors = 1 - cross_validation_accuracies.reshape(-1,1)

    mean_training_accuracies = grid_search.cv_results_['mean_train_score']
    mean_training_errors = 1 - mean_training_accuracies.reshape(-1,1)

    Y_pred = grid_search.best_estimator_.predict(X_test)
    test_accuracy = accuracy_score(Y_pred, Y_test)
    test_error = 1 - sum(Y_pred == Y_test) / len(X_test)

    return opt_K,cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_train
```



KNN for Bank

```
In [94]: knn_bank_accs = []
# 20% Training and 80% Testing
bank_x_train_20, bank_x_test_80, bank_y_train_20, bank_y_test_80 = train_test_split(bank_x, bank_y, test_size=0.2, random_state=0)
opt_K, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors = cross_validation(knn_classifier, bank_x_train_20, bank_y_train_20, bank_x_test_80, bank_y_test_80)

draw_heatmap(cross_validation_errors, k_range, title='cross-validation error w.r.t K')
print("Best C: {}".format(opt_K))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per K: {}".format(mean_training_errors))
print("Avg training accuracies per K: {}".format(mean_training_accuracies))

knn_bank_accs.append(test_accuracy)
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits



Best C: {'n_neighbors': 15}

Test error: 0.11565152478641927

Test accuracy: 0.8843484752135807

Avg training error per K: [[0.]

[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]]


Avg training accuracies per K: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

```
In [95]: # 50% Training and 50% Testing
bank_x_train_50, bank_x_test_50, bank_y_train_50, bank_y_test_50 = train_test_split(bank_x, bank_y, test_size=0.5,
opt_K, cross_validation_accuacies, cross_validation_errors, mean_training_accuacies, mean_training_errors)

draw_heatmap(cross_validation_errors, k_range, title='cross-validation error w.r.t K')
print("Best C: {}".format(opt_K))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per K: {}".format(mean_training_errors))
print("Avg training accuracies per K: {}".format(mean_training_accuacies))

knn_bank_accs.append(test_accuracy)
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits




```
Best C: {'n_neighbors': 14}
Test error: 0.11448288065115453
Test accuracy: 0.8855171193488455
Avg training error per K: [[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]]
Avg training accuracies per K: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```

```
In [96]: bank_x_train_80, bank_x_test_20, bank_y_train_80, bank_y_test_20 = train_test_split(bank_x, bank_y, test_size=0.2, random_state=0)
opt_K, cross_validation_accs, cross_validation_errors, mean_training_accs, mean_training_errors = cross_validate(knn_classifier, bank_x, bank_y, cv=3, scoring='accuracy')

draw_heatmap(cross_validation_errors, k_range, title='cross-validation error w.r.t K')
print("Best C: {}".format(opt_K))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per K: {}".format(mean_training_errors))
print("Avg training accuracies per K: {}".format(mean_training_accs))

knn_bank_accs.append(test_accuracy)
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits



```
Best C: {'n_neighbors': 15}
Test error: 0.10969810903461241
Test accuracy: 0.8903018909653876
Avg training error per K: [[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]]
```

```
Avg training accuracies per K: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```


KNN For INCOME

```
In [97]: knn_income_accs = []
# 20% Training and 80% Testing
income_x_train_20, income_x_test_80, income_y_train_20, income_y_test_80 = train_test_split(income_x, income_y, test_size=0.2, random_state=42)
opt_K, cross_validation_accuacies, cross_validation_errors, mean_training_accuacies, mean_training_errors = cross_validate(KNeighborsClassifier(), income_x_train_20, income_y_train_20, cv=3)

draw_heatmap(cross_validation_errors, k_range, title='cross-validation error w.r.t K')
print("Best C: {}".format(opt_K))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per K: {}".format(mean_training_errors))
print("Avg training accuracies per K: {}".format(mean_training_accuacies))

knn_income_accs.append(test_accuracy)
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits



Best C: {'n_neighbors': 14}

Test error: 0.22934582309582308

Test accuracy: 0.7706541769041769

Avg training error per K: [[0.]

[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]]


Avg training accuracies per K: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

```
In [98]: train_50, income_x_test_50, income_y_train_50, income_y_test_50 = train_test_split(income_x, income_y, test_size=0.5,
cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors, test_acc

ap(cross_validation_errors, k_range, title='cross-validation error w.r.t K')
t C: {}".format(opt_K))
t error: {}".format(test_error))
t accuracy: {}".format(test_accuracy))
training error per K: {}".format(mean_training_errors))
training accuracies per K: {}".format(mean_training_accuracies))

_accs.append(test_accuracy)
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits




```
Best C: {'n_neighbors': 15}
Test error: 0.2147420147420147
Test accuracy: 0.7852579852579853
Avg training error per K: [[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]]
Avg training accuracies per K: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```

```
In [99]: income_x_test_20, income_y_train_80, income_y_test_20 = train_test_split(income_x, income_y, test_size=0.2)
        cross_validation_errors, cross_validation_errors, mean_training_accuracies, mean_training_errors, test_accuracy, test_accuracy = cross_validation_errors, cross_validation_errors, mean_training_accuracies, mean_training_errors, test_accuracy, test_accuracy

        cross_validation_errors, k_range, title='cross-validation error w.r.t K')
        .format(opt_K))
        {}".format(test_error))
        cy: {}".format(test_accuracy))
        g error per K: {}".format(mean_training_errors))
        g accuracies per K: {}".format(mean_training_accuracies))

        pend(test_accuracy)
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits



Best C: {'n_neighbors': 15}

Test error: 0.2071560196560197

Test accuracy: 0.7928439803439803

Avg training error per K: [[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

[1.91945948e-05]

Avg training accuracies per K: [0.99998081 0.99998081 0.99998081 0.99998081 0.99998081 0.99998081

0.99998081 0.99998081 0.99998081 0.99998081 0.99998081 0.99998081

0.99998081 0.99998081 0.99998081]

KNN For MUSHROOM


```
In [100]: knn_mushroom_accs = []

# 20% Training and 80% Testing
mushroom_x_train_20, mushroom_x_test_80, mushroom_y_train_20, mushroom_y_test_80 = train_test_split(income,
opt_K, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, k_range, title='cross-validation error w.r.t K')
print("Best K: {}".format(opt_K))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per K: {}".format(mean_training_errors))
print("Avg training accuracies per K: {}".format(mean_training_accuracies))

knn_mushroom_accs.append(test_accuracy)
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits



Best K: {'n_neighbors': 15}

Test error: 0.2264665233415234

Test accuracy: 0.7735334766584766

Avg training error per K: [[0.]

[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]
[0.]]

Avg training accuracies per K: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]

In [102]: # 50% Training and 50% Testing

```


mushroom_x_train_50, mushroom_x_test_50, mushroom_y_train_50, mushroom_y_test_50 = train_test_split(income,
opt_K, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, k_range, title='cross-validation error w.r.t K')
print("Best K: {}".format(opt_K))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per K: {}".format(mean_training_errors))
print("Avg training accuracies per K: {}".format(mean_training_accuracies))

cnn_mushroom_accs.append(test_accuracy)

```

Fitting 3 folds for each of 15 candidates, totalling 45 fits



Best K: {'n_neighbors': 15}

Test error: 0.21517199017199018

Test accuracy: 0.7848280098280098

Avg training error per K: $[[6.14241183e-05]$

```
[6.14241183e-05]
```

[6.14241183e-05]

```
[6.14241183e-05]
```

[6.14241183e-05]

[6.14241183e-05]

[6.14241183e-05]

```
[6.14241183e-05]
```

[6.14241183e-05]

[6.14241183e-05]

```
[6.14241183e-05]
```

[6.14241183e-05]

[6.14241183e-05]

```
[6.14241183e-05]
```

```
[6.14241183e-05]]
Avg training accuracies per K: [0.99993858 0.99993858 0.99993858 0.99993858 0.99993858 0.99993858
```


0.99993858 0.99993858 0.99993858 0.99993858 0.99993858 0.99993858

```
0.99993858 0.99993858 0.99993858]
```



```
In [103]: testing
mushroom_x_test_20, mushroom_y_train_80, mushroom_y_test_20 = train_test_split(income_x, income_y, test_size=0.2, random_state=42)
accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors, test_accuracy, test_error = cross_validate(knn_classifier, mushroom_x_train_80, mushroom_y_train_80, cv=3, scoring='accuracy')
cross_validation_errors, k_range, title='cross-validation error w.r.t K')
plt.plot(k_range, cross_validation_errors)
plt.xlabel('K')
plt.ylabel('cross-validation error')
plt.title(title)
plt.show()
print('Test error: %.4f' % test_error)
print('Test accuracy: %.4f' % test_accuracy)
print('Avg training error per K: %.4f' % np.mean(mean_training_errors))
print('Avg training accuracies per K: %.4f' % np.mean(mean_training_accuracies))
```

Fitting 3 folds for each of 15 candidates, totalling 45 fits



```
Best K: {'n_neighbors': 15}
Test error: 0.20915233415233414
Test accuracy: 0.7908476658476659
Avg training error per K: [[3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]
 [3.8390295e-05]]
Avg training accuracies per K: [0.99996161 0.99996161 0.99996161 0.99996161 0.99996161 0.99996161
 0.99996161 0.99996161 0.99996161 0.99996161 0.99996161 0.99996161 0.99996161 0.99996161
 0.99996161 0.99996161 0.99996161]
```

```
In [104]: average_knn_accuracy = np.sum([a + b + c for a, b, c in zip(knn_bank_accs, knn_income_accs, knn_mushroom_accs)])
print(average_knn_accuracy)
0.8175703089297898
```

Decision Tree

```
In [106]: import seaborn as sns
from sklearn import tree

D_list = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

```
In [107]: def calcDTMetrics(X_train,X_test, Y_train,Y_test, D_list):

    estimator = tree.DecisionTreeClassifier(criterion='entropy',random_state = 1 )

    param_grid = {'max_depth': D_list}
    grid_search = GridSearchCV(estimator, param_grid, cv=3, return_train_score=True)

    # Fit the model
    grid_search.fit(X_train, Y_train)

    # Gather the results
    opt_D = grid_search.best_params_['max_depth']

    cross_validation_accuracies = grid_search.cv_results_['mean_test_score']
    cross_validation_errors = 1 - cross_validation_accuracies.reshape(-1,1)

    mean_training_accuracies = grid_search.cv_results_['mean_train_score']
    mean_training_errors = 1 - mean_training_accuracies.reshape(-1,1)

    Y_pred = grid_search.best_estimator_.predict(X_test)
    test_accuracy = accuracy_score(Y_pred, Y_test)
    test_error = 1 - sum(Y_pred == Y_test) / len(X_test)


    return opt_D ,cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_train
```

Decision Tree for BANK

```
In [128]: dt_bank_accs = []
# 20% Training and 80% Testing
bank_x_train_20, bank_x_test_80, bank_y_train_20, bank_y_test_80 = train_test_split(income_x, income_y, test_size=0.2, random_state=42)
opt_D, cross_validation_accs, cross_validation_errors, mean_training_accs, mean_training_errors = cross_validate(DT, bank_x_train_20, bank_y_train_20, cv=5)

draw_heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
print("Best D: {}".format(opt_D))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per D: {}".format(mean_training_errors))
print("Avg training accuracies per D: {}".format(mean_training_accs))

dt_bank_accs.append(test_accuracy)
```



Best D: 5

Test error: 0.15582770270270274

Test accuracy: 0.8441722972972973

Avg training error per D: [[0.24646806]

[0.18542694]

[0.16369683]

[0.16369683]

[0.15041451]

[0.14534714]

[0.13951186]

[0.13075888]

[0.12231271]


[0.11425062]]

Avg training accuracies per D: [0.75353194 0.81457306 0.83630317 0.83630317 0.84958549 0.85465286 0.86048814 0.86924112 0.87768729 0.88574938]

```
In [129]: # 50% Training and 50% Testing
bank_x_train_50, bank_x_test_50, bank_y_train_50, bank_y_test_50 = train_test_split(income_x, income_y, test_size=0.5, random_state=42)
opt_D, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors = grid_search(D_list)

draw_heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
print("Best D: {}".format(opt_D))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per D: {}".format(mean_training_errors))
print("Avg training accuracies per D: {}".format(mean_training_accuracies))

dt_bank_accs.append(test_accuracy)
```




```
Best D: 7
Test error: 0.14533169533169532
Test accuracy: 0.8546683046683047
Avg training error per D: [[0.24146192]
 [0.17195945]
 [0.15795454]
 [0.15795454]
 [0.15540538]
 [0.14763514]
 [0.14345823]
 [0.13799139]
 [0.13175673]
 [0.12610566]]
Avg training accuracies per D: [0.75853808 0.82804055 0.84204546 0.84204546 0.84459462 0.85236486
 0.85654177 0.86200861 0.86824327 0.87389434]
```

```
In [130]: # 80% Training and 20% Testing
bank_x_train_80, bank_x_test_20, bank_y_train_80, bank_y_test_20 = train_test_split(income_x, income_y, test_size=0.2, random_state=42)
opt_D, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors = grid_search()

draw_heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
print("Best D: {}".format(opt_D))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per D: {}".format(mean_training_errors))
print("Avg training accuracies per D: {}".format(mean_training_accuracies))

dt_bank_accs.append(test_accuracy)
```




```
Best D: 8
Test error: 0.14296683046683045
Test accuracy: 0.8570331695331695
Avg training error per D: [[0.24105498]
 [0.18673203]
 [0.1561156 ]
 [0.15600043]
 [0.15189261]
 [0.14630682]
 [0.14294764]
 [0.13918534]
 [0.13484716]
 [0.13087369]]
Avg training accuracies per D: [0.75894502 0.81326797 0.8438844  0.84399957 0.84810739 0.85369318
 0.85705236 0.86081466 0.86515284 0.86912631]
```

Decision Tree for INCOME

```
In [114]: dt_income_accs = []
# 20% Training and 80% Testing
income_x_train_20, income_x_test_80, income_y_train_20, income_y_test_80 = train_test_split(income_x, income_y,
                                                                                               test_size=0.2,
                                                                                               random_state=0)
opt_D, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors = cross_validate(
    DecisionTreeClassifier(), income_x_train_20, income_y_train_20, cv=5, n_jobs=-1)

draw_heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
print("Best D: {}".format(opt_D))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per D: {}".format(mean_training_errors))
print("Avg training accuracies per D: {}".format(mean_training_accuracies))

dt_income_accs.append(test_accuracy)
```



Best D: 6

Test error: 0.1491477272727273

Test accuracy: 0.8508522727272727

Avg training error per D: [[0.24017199]

[0.17813322]

[0.15555905]

[0.15356258]

[0.14933948]

[0.14419494]

[0.13651711]

[0.1305282]

[0.12215894]


[0.11417378]]

Avg training accuracies per D: [0.75982801 0.82186678 0.84444095 0.84643742 0.85066052 0.85580506
0.86348289 0.8694718 0.87784106 0.88582622]

```
In [115]: # 50% Training and 50% Testing
income_x_train_50, income_x_test_50, income_y_train_50, income_y_test_50 = train_test_split(income_x, income_y,
opt_D, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
print("Best D: {}".format(opt_D))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per D: {}".format(mean_training_errors))
print("Avg training accuracies per D: {}".format(mean_training_accuracies))

dt_income_accs.append(test_accuracy)
```




```
Best D: 7
Test error: 0.14502457002457003
Test accuracy: 0.85497542997543
Avg training error per D: [[0.24183047]
 [0.19582305]
 [0.1579238 ]
 [0.15740168]
 [0.14944703]
 [0.14376527]
 [0.14078622]
 [0.13713144]
 [0.13280101]
 [0.12616711]]
Avg training accuracies per D: [0.75816953 0.80417695 0.8420762 0.84259832 0.85055297 0.85623473
0.85921378 0.86286856 0.86719899 0.87383289]
```

In [116]: *Training and 20% Testing*

```
income_x_train_80, income_x_test_20, income_y_train_80, income_y_test_20 = train_test_split(income_x, income_y, test_size=0.2, random_state=42)
cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors, test_accuracy = cross_validate(model, income_x_train_80, income_y_train_80, income_x_test_20, income_y_test_20, cv=10)

heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
"Best D: {}".format(opt_D)
"Test error: {}".format(test_error)
"Test accuracy: {}".format(test_accuracy)
"Avg training error per D: {}".format(mean_training_errors)
"Avg training accuracies per D: {}".format(mean_training_accuracies)

ome_accs.append(test_accuracy)
```



Best D: 8

Test error: 0.14388820638820643

Test accuracy: 0.8561117936117936

Avg training error per D: [[0.24078624]

[0.18629052]

[0.15644192]

[0.15626916]

[0.15256445]

[0.14590369]

[0.14175751]

[0.13724662]

[0.13331161]

[0.12922302]]

Avg training accuracies per D: [0.75921376 0.81370948 0.84355808 0.84373084 0.84743555 0.85409631 0.85824249 0.86275338 0.86668839 0.87077698]


Decision Tree for Mushroom

```
In [117]: dt_mushroom_accs = []

# 20% Training and 80% Testing
mushroom_x_train_20, mushroom_x_test_80, mushroom_y_train_20, mushroom_y_test_80 = train_test_split(income,
opt_D, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
print("Best D: {}".format(opt_D))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per D: {}".format(mean_training_errors))
print("Avg training accuracies per D: {}".format(mean_training_accuracies))

dt_mushroom_accs.append(test_accuracy)
```




```
Best D: 6
Test error: 0.1515663390663391
Test accuracy: 0.8484336609336609
Avg training error per D: [[0.23141892]
[0.18481252]
[0.15164315]
[0.15018428]
[0.14189218]
[0.13928161]
[0.13490495]
[0.13022131]
[0.12238957]
[0.11463453]]
Avg training accuracies per D: [0.76858108 0.81518748 0.84835685 0.84981572 0.85810782 0.86071839
0.86509505 0.86977869 0.87761043 0.88536547]
```

```
In [118]: # 50% Training and 50% Testing
mushroom_x_train_50, mushroom_x_test_50, mushroom_y_train_50, mushroom_y_test_50 = train_test_split(income,
opt_D, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
print("Best D: {}".format(opt_D))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per D: {}".format(mean_training_errors))
print("Avg training accuracies per D: {}".format(mean_training_accuracies))

dt_mushroom_accs.append(test_accuracy)
```




```
Best D: 8
Test error: 0.149017199017199
Test accuracy: 0.850982800982801
Avg training error per D: [[0.23937346]
[0.17801008]
[0.15605039]
[0.15595826]
[0.15205775]
[0.14404186]
[0.13977278]
[0.13445947]
[0.12954551]
[0.12361802]]
Avg training accuracies per D: [0.76062654 0.82198992 0.84394961 0.84404174 0.84794225 0.85595814
0.86022722 0.86554053 0.87045449 0.87638198]
```

```
In [119]: # 80% Training and 20% Testing
mushroom_x_train_80, mushroom_x_test_20, mushroom_y_train_80, mushroom_y_test_20 = train_test_split(income,
opt_D, cross_validation_accuracies, cross_validation_errors, mean_training_accuracies, mean_training_errors)

draw_heatmap(cross_validation_errors, D_list, title='cross-validation error w.r.t D')
print("Best D: {}".format(opt_D))
print("Test error: {}".format(test_error))
print("Test accuracy: {}".format(test_accuracy))
print("Avg training error per D: {}".format(mean_training_errors))
print("Avg training accuracies per D: {}".format(mean_training_accuracies))

dt_mushroom_accs.append(test_accuracy)
```



```
Best D: 10
Test error: 0.14619164619164615
Test accuracy: 0.8538083538083538
Avg training error per D: [[0.24132371]
 [0.19560043]
 [0.15694102]
 [0.15669149]
 [0.15091377]
 [0.14640284]
 [0.14171914]
 [0.13849433]
 [0.13363791]
 [0.12828241]]
Avg training accuracies per D: [0.75867629 0.80439957 0.84305898 0.84330851 0.84908623 0.85359716
 0.85828086 0.86150567 0.86636209 0.87171759]
```

```
In [120]: average_dt_accuracy = np.sum([a + b + c for a, b, c in zip(dt_bank_accs, dt_income_accs, dt_mushroom_accs)])
print(average_dt_accuracy)

0.8509836541086542
```

In [132]:

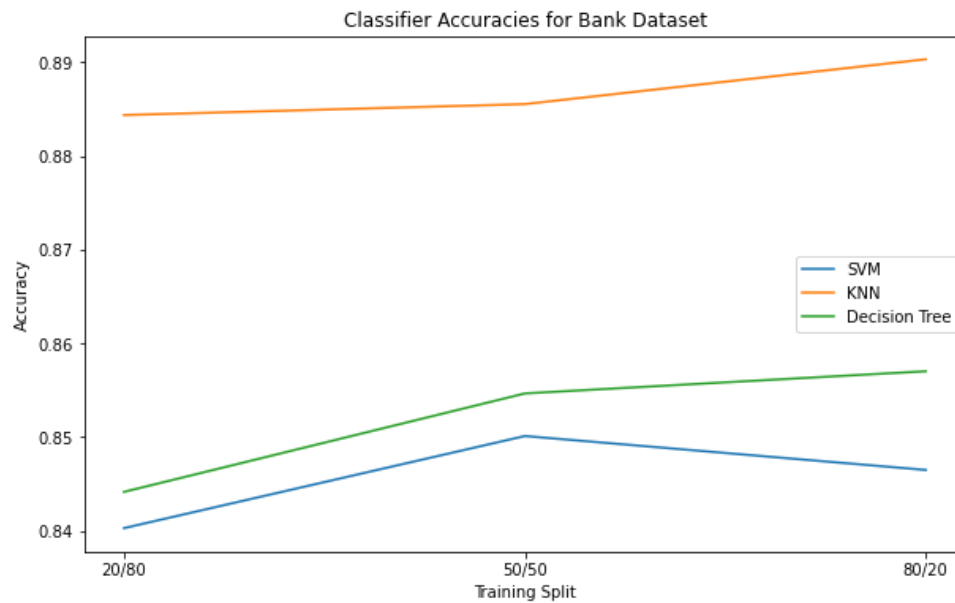
```
data = {
    'SVM': [svm_bank_accs, svm_income_accs, svm_mushroom_accs],
    'KNN': [knn_bank_accs, knn_income_accs, knn_mushroom_accs],
    'Decision Tree': [dt_bank_accs, dt_income_accs, dt_mushroom_accs],
}

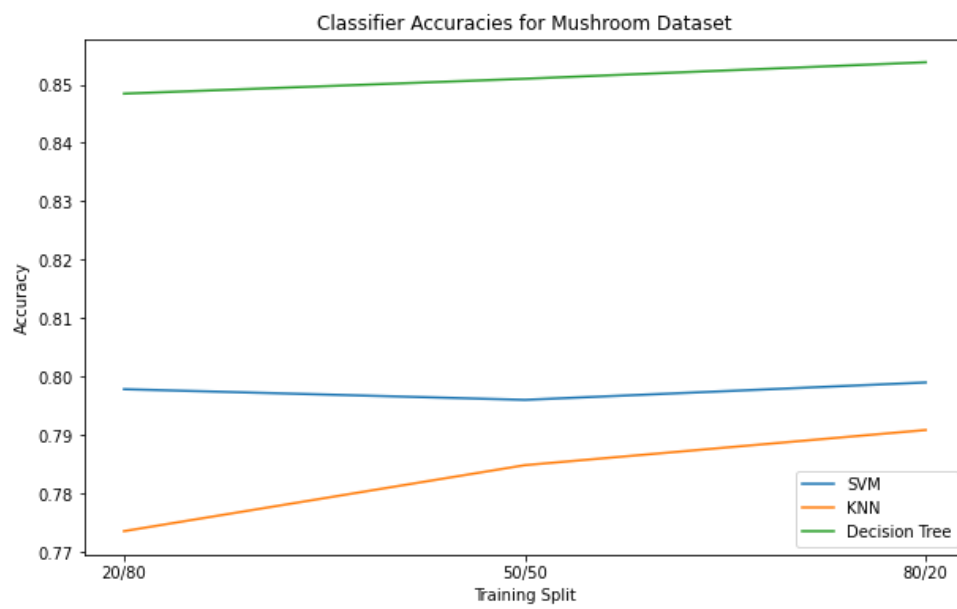
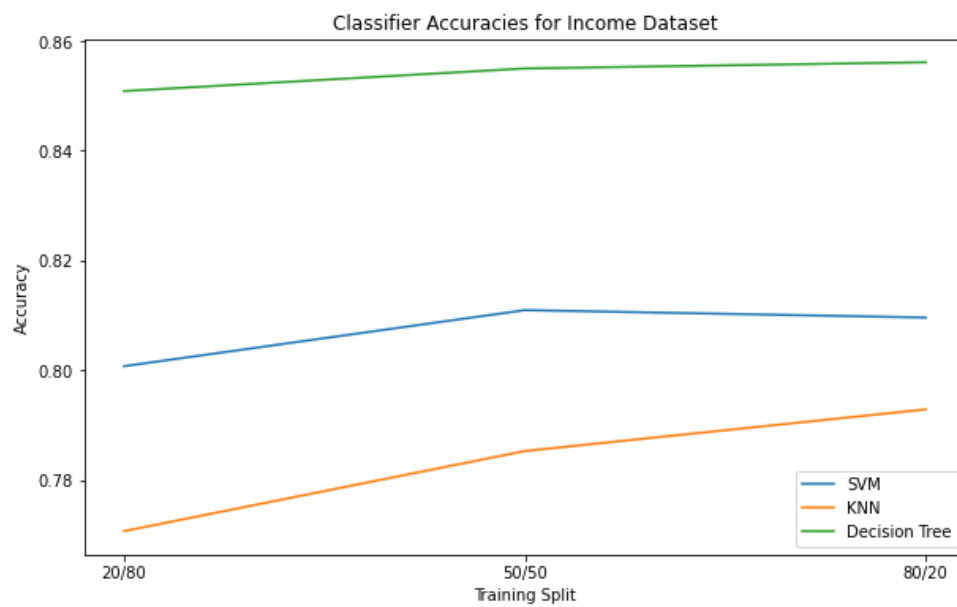
# Data preparation
datasets = ['Bank', 'Income', 'Mushroom']
classifiers = ['SVM', 'KNN', 'Decision Tree']

# Plotting
for i, dataset in enumerate(datasets):
    plt.figure(figsize=(10, 6))
    plt.title(f'Classifier Accuracies for {dataset} Dataset')
    plt.xlabel('Training Split')
    plt.ylabel('Accuracy')

    for classifier in classifiers:
        accs = data[classifier][i]
        splits = ['20/80', '50/50', '80/20']
        plt.plot(splits, accs, label=classifier)

plt.legend()
plt.show()
```





```
In [163]: bar_width = 0.2 # Width of each bar
bar_positions = np.arange(len(datasets))
colors = plt.cm.Set3(np.linspace(0, 1, len(classifiers)))

for i, split in enumerate(['20/80', '50/50', '80/20']):
    plt.figure(figsize=(10, 6))
    plt.title(f'Classifier Accuracies for {split} Training Split')
    plt.xlabel('Dataset')
    plt.ylabel('Accuracy')
    plt.xticks(bar_positions, datasets)

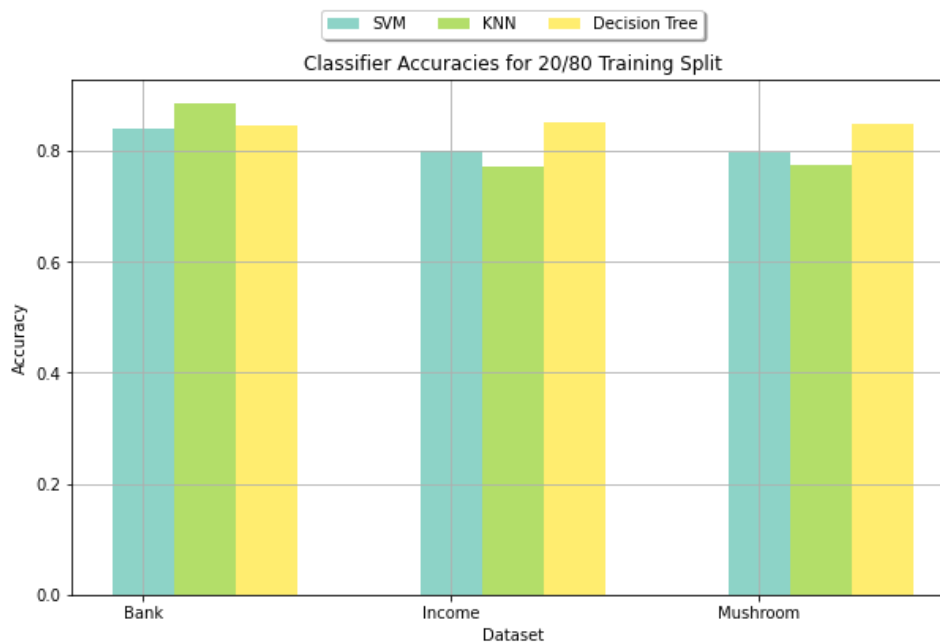
    for j, (classifier, color) in enumerate(zip(classifiers, colors)):
        accs = [data[classifier][k][i] for k in range(len(datasets))]
        plt.bar(bar_positions + j * bar_width, accs, width=bar_width, label=classifier, color=color)

        avg_acc = np.mean(accs)
        print("Average Accuracy for {}: {}".format(classifier, avg_acc))

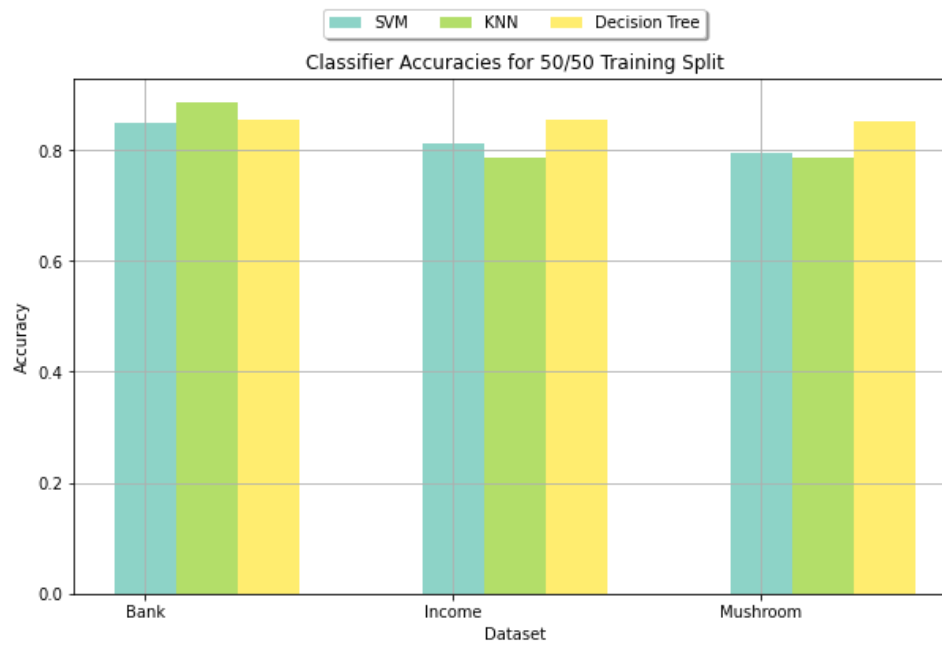
    plt.legend(loc='upper center', bbox_to_anchor=(0.5, 1.15), fancybox=True, shadow=True, ncol=len(classifiers))

    plt.grid(True)
    plt.show()
```


Average Accuracy for SVM: 0.8129513555625177
Average Accuracy for KNN: 0.8095120429254115
Average Accuracy for Decision Tree: 0.8478194103194102



Average Accuracy for SVM: 0.8190231054890879
Average Accuracy for KNN: 0.8185343714782801
Average Accuracy for Decision Tree: 0.8535421785421785



Average Accuracy for SVM: 0.8183599698791006
Average Accuracy for KNN: 0.8246645123856778
Average Accuracy for Decision Tree: 0.8556511056511056




```
In [123]: average_accuracies = [np.mean(np.array(data[classifier])) for classifier in classifiers]

df_data = pd.DataFrame({
    'Classifier': classifiers,
    'Average Accuracy': average_accuracies
})

# Create a horizontal bar plot
plt.figure(figsize=(10, 6))
sns.barplot(x='Average Accuracy', y='Classifier', data=df_data, palette='viridis', ci=None)
plt.title('Average Accuracies of Classifiers')
plt.xlabel('Average Accuracy')
plt.ylabel('Classifier')

# Add Labels for each point
for i, v in enumerate(average_accuracies):
    plt.text(v + 0.01, i, f'{v:.3f}', color='black', va='center', fontweight='bold')

plt.show()
```





```
In [124]: # Convert accuracy data to a 2D NumPy array
heatmap_data = np.array([[data[classifier][i][j] for j in range(len(datasets))] for i in range(len(datasets))])

# Plotting the heatmap
plt.figure(figsize=(10, 6))
plt.imshow(heatmap_data, cmap='viridis', interpolation='nearest', aspect='auto')

plt.colorbar(label='Accuracy')
plt.title('Classifier Accuracies Heatmap')
plt.xlabel('Dataset')
plt.ylabel('Classifier')
plt.xticks(np.arange(len(datasets)), datasets)
plt.yticks(np.arange(len(classifiers)), classifiers)

plt.show()
```



In []: